

Systems Programming

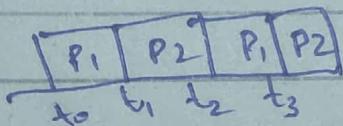
Book: Unix Systems Programming: Communication
Concurrency & Threads Author K. A. Robbins &
Steven Robbins

Contents :

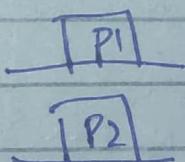
- Midterm

 1. _____ self study → Reading Assignment
 2. Program, Process, threads
 3. Processes
 4. Unix I/O
(standard I/O device)
(Regular files) represent memory address of function
main() {
 scanf(); → sys call
 5. Files & Directories
 6. Special files (fifos / pipes / sockets)
 7. Signals

Concurrency

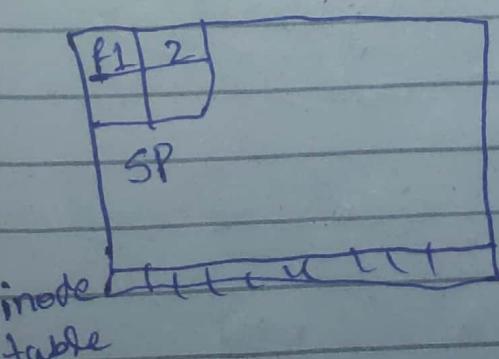


Parallelism



- q. Times & Timers
 - 13. Threads / 15. IPC using shared Memory.

400



Signal:- notification
of event.

CLO's :-

CLO1: error handling / file handling /
signal handling.

F.Term

CLO2: parallel processing (multiprocessing /
multithreading)

F.Term CLO3: Interprocess Communication for
real world.

Grading Criteria:-

Mid: 30%

Final: 40%

Sessional: 30%

Project / Attendance / Quizzes / Assignment

(3)

(3)

Oct 5, 2023

Programs, Process & Threads

Program → Process

#include < _____ > } → pre-compiler, preprocessor
#define COUNT 30 } → directives, directives
→ do not compile,

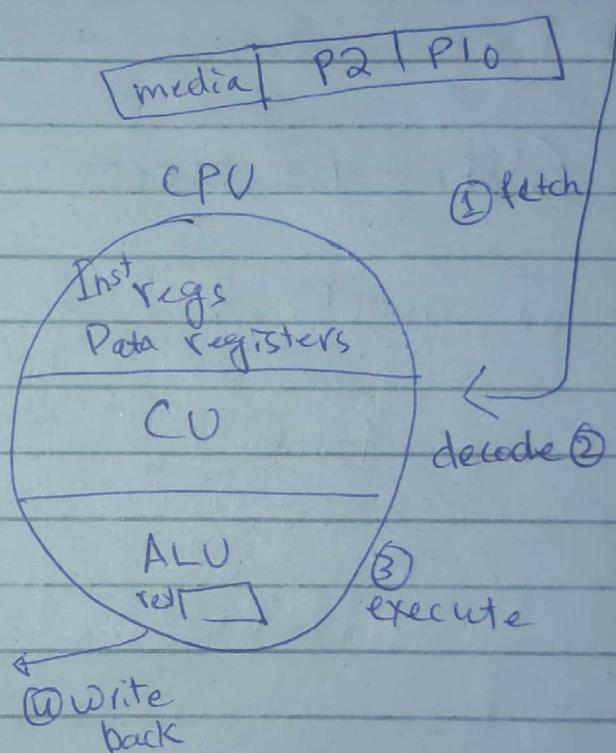
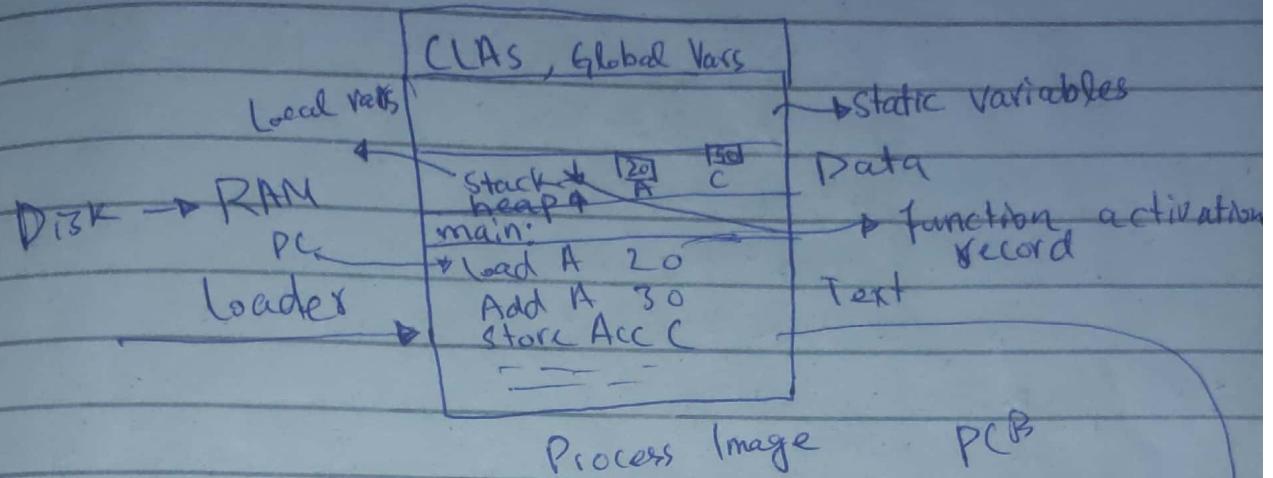
int main(int argc, char * argv[]){
 int A = 20;

int C = A + COUNT

printf("%d\n", C);

return(0); → exit status returned to parent

gcc P1.C -o P1.O
 disk disk



Variables
 Scope
 Usage
 lifetime

```
int main() {
    int A = 20; int C = A + COUNT;
    fun(A); → Stack will be
    divided
```

function stack will be destroyed when it returns.

Fun(int A){

int X = 30;

int Y = 40;

X++; Y++;

cout << X << Y;

}

* Dynamic Memory must be deallocated

Static Variables:

```
int main() {  
    func();  
    func();  
    return 0;  
}
```

void fun() {

int A = 10;

static int B = 20;

dec + init [execute only once]

A++;

B++;

cout << A << B;

Compile time,
part of exec mode

Compile
not part
of exec
module

CLL's, Gr. Variables

initialized

uninitialized

Stack ↓

Heap ↑

Ketna

PCB

A linked list.

They

Local vars are destroyed while static aren't.

```
void fun() {  
    static int c;  
    c = 10;  
    c++;  
    cout << c;
```

Static int A [100];

400 bytes

static int A [100] =
{1, 2, 3};

$$400 + 100 * 4 + 7 \\ = 807$$

ste
one

Oct 6, 2023

Library function Calls:

PID
Process State
PC
Registers
Statistics
Memory limits
List of opened files

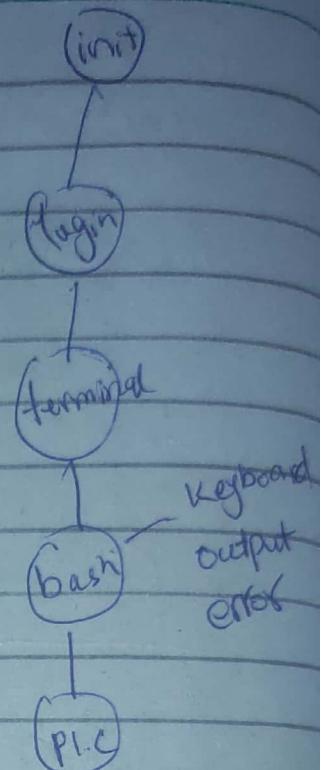
file descriptor table has 3

Keyboards	0	STDIN_FILENO	files by default
	1	STDOUT_FILENO	→ monitor
	2	STDERR_FILENO	→ monitor

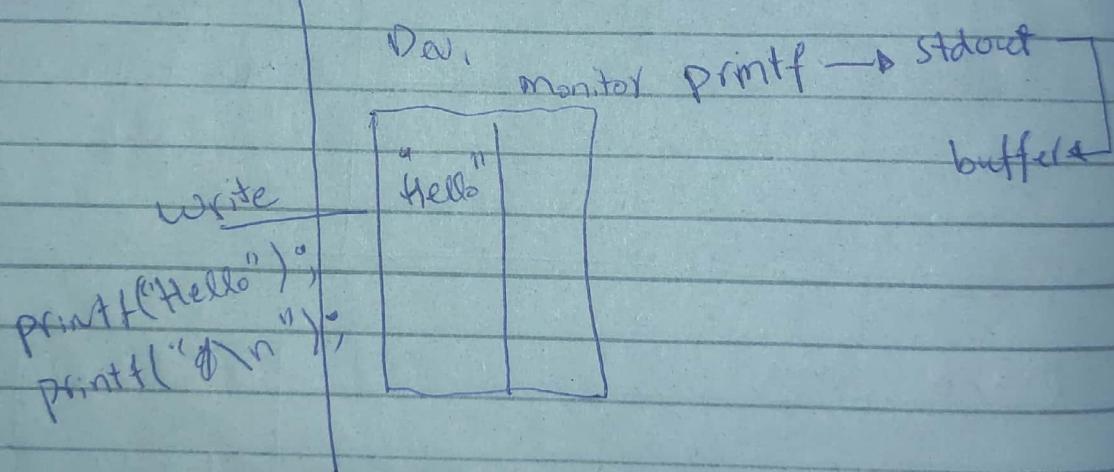
→ Child process will have
Same PC, registers and
list of opened files.

stdin, stdout, stderr are
opened by default.

* Why two entries for
monitor.



Stdout | vs Stderror



- | | |
|-----------------------|-----------------|
| (1) new line | (3) fflush |
| (2) input | (4) buffer full |
| (5) process terminate | |

→ We deal with device buffer in order
to save cycles.

→ In error case, we want to
let user know immediately.
So stderror has no buffer.

→ buffer is memory in RAM.

```
main() {  
    printf("Hello1"); → stdout  
    printf ("Hello2"); → stderr  
}
```

Output: Hello2 Success
Hello1

filename string
fprintf(stdout, "Hello1");
fprintf(stderr, "Hello2");

Library function calls :-

- ① Detect error
- ② Error Reason

return-type name(args);

int

int close(int fd);

-1 : error
0 : success

0	Std in
1	Std out
2	error

```

errno
↑
errno.h
list of errors
0
1
2
3
8 Invalid fd

```

```

int main()
{
    int ret = close(4);
    if (ret == -1) { // error
        perror("error");
        fprintf(stderr, "Failed to close\n"
                ": %s", strerror(errno));
    }
    return -1;
}

```

② set errno
-1 (0 Retention)
failure

8th
char * strerror(int errno);

(close()) → will close std::out
and printf will fail.

Date 9 Oct, 2023

Process Environment

→ \$ env

↳ List Environment
Vars

Data

ls → pwd



PATH



$\text{HOME} = \text{/home/student}$

echo \$HOME

Environment vars are also inherited by child process.

extern char ** environ

↳ Let compiler know that we are using it from another source file.

→ it will not be declared as a new var.

main() {

```
for( int i=0; environ[i] !=NULL; i++)
    printf("%\n", environ[i]);
}
```

char * getenv(const char * varname);

main() {

char * valofpath;

valofpath = getenv("PATH");

printf ("PATH=%s\n", valofpath);

Normal Process Termination:

- Exit handlers
- stty
1. Explicit `exit` → different
 2. Implicit exit (return statement of main)
 3. `-exit()` / `-Exit()` Same

`int atexit(void * fun(void))`

0: success

non-zero: failure

```
int main() {
    fun1
    atexit('fun1');
    atexit('fun2');
    ① printf("In main\n");
    exit(0);
}
```

`-exit(0); // does not call
// exit handlers`

Note:
`fun1(void) {`
② `printf("Process
Terminating Normally
\n");`
}

`void fun2(void) {`
③ `printf("All sub
executed\n");`
}

```

int main () {
    atexit( fun1 );
    atexit( fun2 );
    int ret = close(4);
    if (ret == -1) {
        fprintf( stderr, " Failed close sc" );
        exit(-1);
    }
    return 0;
}

```

→ String Tokenization Skipped.

→ Chap 2 Quiz in next class.

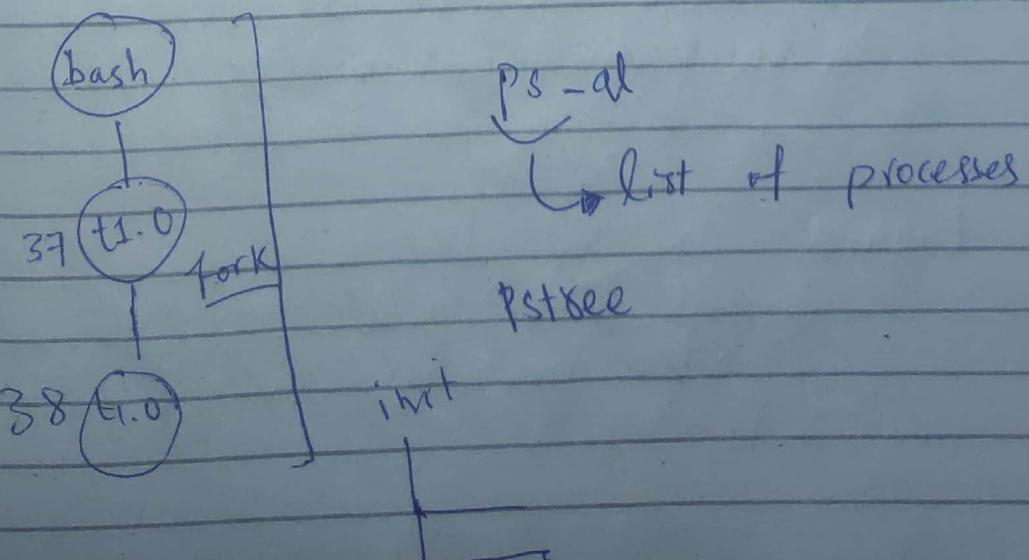
Chapter (3)

SP

16/10/2023

Processes:

PCB P3	PCB P2	PCB P1
PID PC State		



for making a group,
 sudo groupadd Section C
 sudo usermod -aG SectionC Hamza
 Append 3 14

1st Column	2nd Column	state
ID	D	uninterruptible sleep
	S	interruptible sleep, wait
	R	
	Z	Zombie
	T	Signals

```
int getpid();
int getppid();
int getuid();
int getgid();
```

root root
0 0

Effective ID's:-

user: int geteuid();
 group: int getegid();

Logged in as Hamza

Sudo groupadd myfriends

Let's test it:

```
main() {
```

```
    printf (" My ID = %d , My GID = %d \n",  
           getuid(), getgid());
```

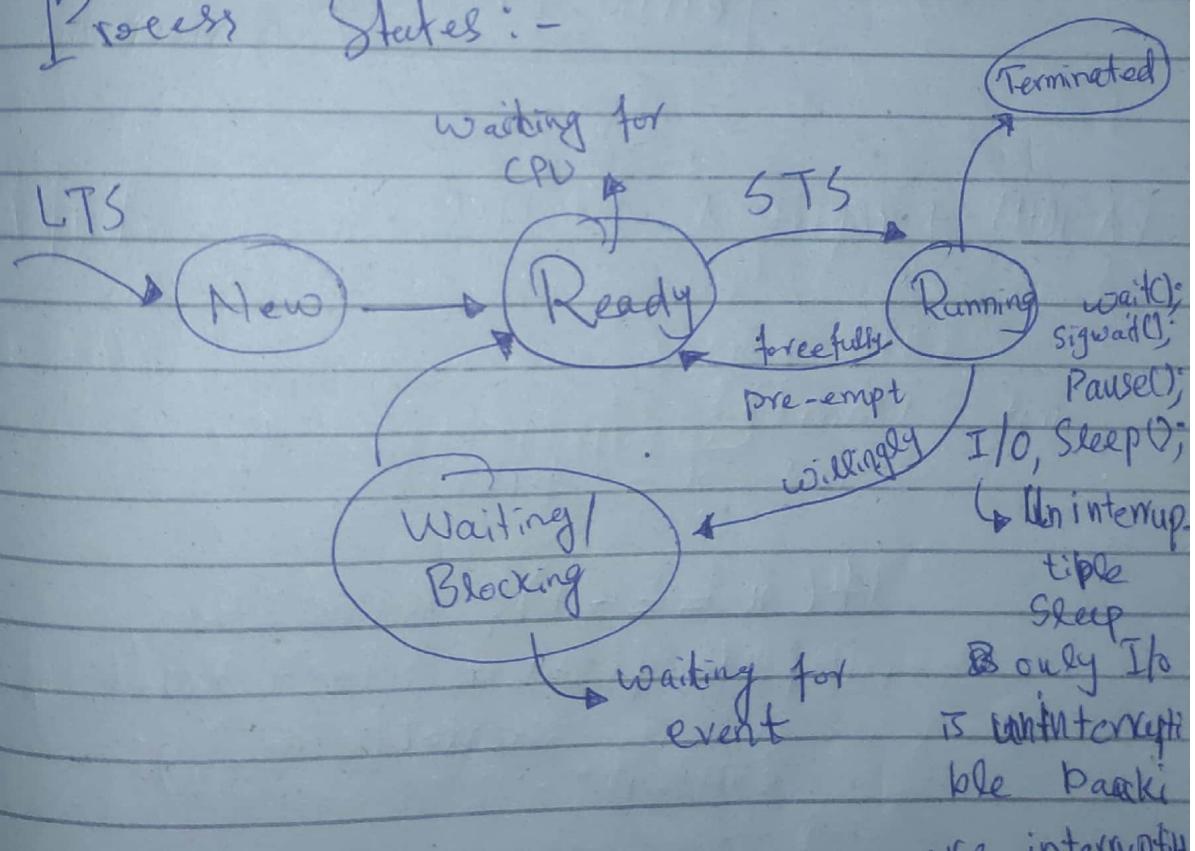
```
}
```

sudo ./t10 → My ID = 14, GID = 4
My ID = 0, GID = 0

Effective ID's

↳ Temporarily changed

Process States :-



When child process terminate, parent receive SIGCHLD signal.

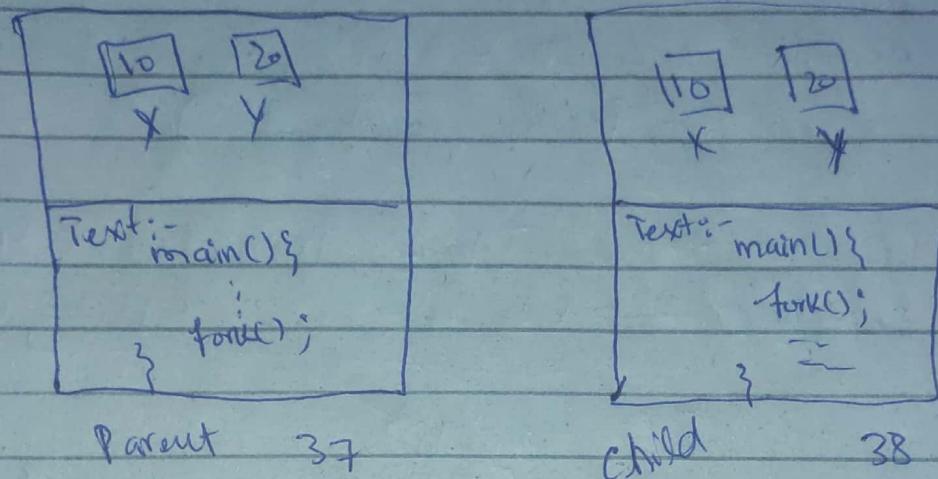
RD: 50
PL: 0x50
State
fd table

PC: 0x50
State
fd table

(1) no resources
(2) limited
↑ ↑ ↑

Multiprocessing:

int fork();
-1 : error
0 :
>0 :



→ List of all errors [HW]

errno: EAGAIN

print int

↳ corresponding
String

EAGAIN → represent the error.

```
int main(){
    int x = fork();
    if(x == -1){ // failure
        fprintf(stderr, "fork Failed o/p d\n",
                strerror(errno));
        exit(-1);
    }
    printf("Im o/p d , my x = o/p d\n",
           getpid(), x);
    return 0;
}
```

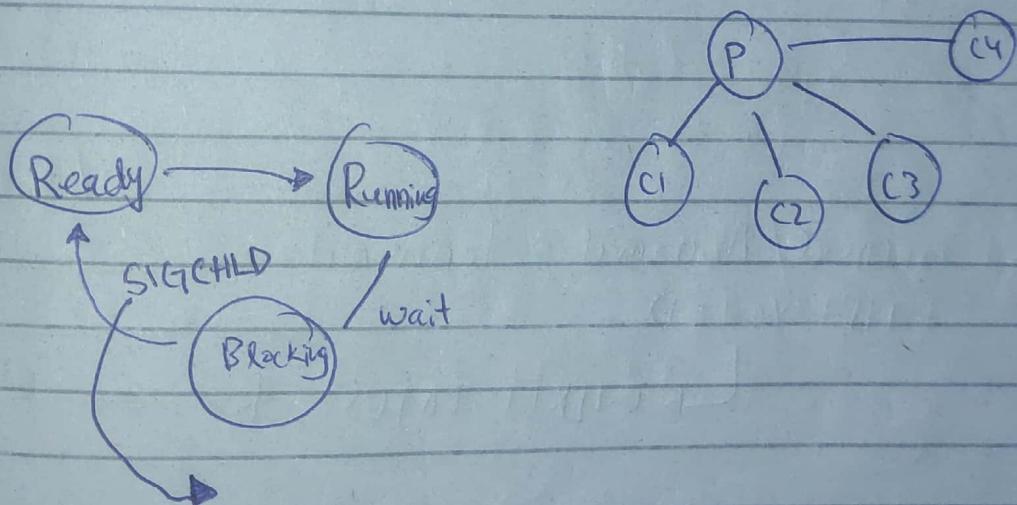
Date 19 Oct 2023

Processes:-

int fork()
error: -1, errno
Success: 0

> 0 : parent
child ID

int wait (int *status) → child reasons of termination.
error: -1, errno
Success: > 0 no child exist.
child ID ① ECHILD
② EINTR
③ EINVAL
only in waitpid



int waitpid (int pid, int *s, int option)

Process Termination

Normal → Exit Status returned to parent in wait argument.

exit
return
- exit
- Exit

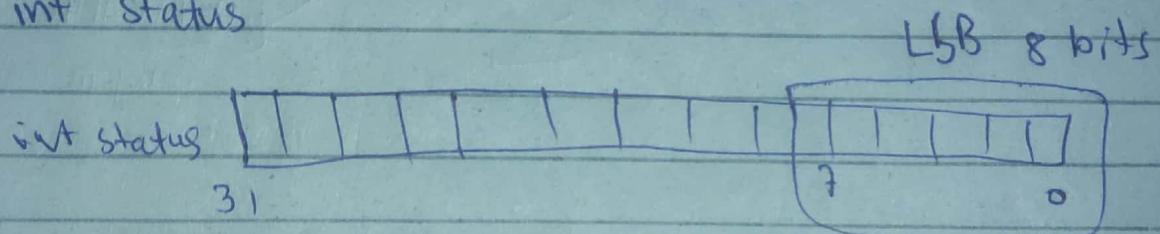
ABNORMAL → terminated completely
terminated → suspend → stopping sig

```

int main() {
    int x = fork();
    if (x == 0) {
        printf("I'm child");
        exit(4);
    }
    int status;
    int ret = wait(&status);
}

```

int status



For normal/Abnormal termination,
WIFEXITED

↳ WEXITSTATUS for ex value

Now for abnormal termination,
WIFSIGNALED

↳ for find terminating signal

WTERMSIGNAL

For last scenario; WIFSTOPPED

↳ WSTOPSIG

```
→ if (WIFEXITED(status))
    printf("Child exited normally with
           Status = %d\n", WEXITSTATUS(status));
if (WIFSIGNALED(status))
    printf("Terminating Signal = %d\n",
           WTERMSIG(status));
if (WIFSTOPPED(status))
    printf("Stoping Signal = %d\n", WSTOPSIG(
           status));
```

```
int main() {
    int N = 20;
    int cid[N];
    for (int i=0; i<N; i++) {
        cid[i] = fork();
        if (cid[i] == 0) {
            printf("I'm child");
            exit(0);
        }
        else if (cid[i] == -1) {
            printf(stderr, "Fork failed: %s\n",
                   strerror(errno));
        }
    }
    return 1;
}
```

int status
⑧

int ret = waitpid(CMD[3], &status, 1);

-exit(4) → does not empty std::cout buffer.

printf("Before for loop\n"); will be

Printed 6 times.

Date 26/10/2023

Execute System Call:

int exec, int execp
int execv, int execvp
int main() {
 int execle, int execlve
 execp("ls", "ls", NULL);
 printf("Hello");
}

→ Path

→ will not print

error -1

Success: → will not return.

errno: E2BIG

EINVAL

EACCES

|

Contains built-in
commands Path → Searches
in PATH

int main → Complete Path → from which
exec ("bin/ls", command.

PATH = /bin: /usr/bin: /lib: /usr/lib
① ② ③ ④ ⑤

which t1.o → Command not found.

int exec (const char *path, const char *filename,
 const char *arg0, const char *arg1, ---
 ---, const char *argN, NULL);

int execvp (const char *filename, const char *
 arg0, ---, const char *argN, NULL);

* / t1.o 1 2

int main (int argc, char *argv[7]) {
 if (argc < 3) {
 printf ("Invalid Args");
 return -1;
 }
 + sum;

```
Sum = atoi(argv[1]) + atoi(argv[2]);
```

```
printf("I'm %s, sum = %d\n", argv[0],  
      sum);  
return sum;  
}
```

t2.c

```
int main() {  
    int argc, char *argv[];  
    . /t2.o  
  
    int x = fork();  
    if (x == 0) // child  
        execl("./t1.o", "t1.o", "1", "2",  
              argv[1], argv[2]);  
    else  
        int s = wait(&s);  
        printf("Child process returned sum  
              = %d\n", WEXITSTATUS(s));  
}
```

```
int execv(const char* filenamePath, const  
          @ char* argv[])
```

```
int x = fork();  
if (x == 0) // child  
{  
    execvp("./t1.o", &argv[1]);  
    address of element
```

t2. c

```
int main( int argc, char* argv[] )  
int x = fork()
```

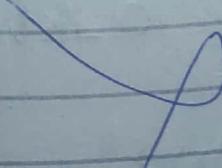
t2. c

```
int main( int argc, char* argv[] )
```

```
char *newenv[] = { "VAR1=Batch23",  
                   "VAR2=Sect.C",  
                   NULL };
```

```
int x = fork();  
if (x == 0) {  
    // print old env.  
    execve("./t1.0", &argv[1], newenv);  
}  
wait(NULL);  
}
```

Date 30th Oct, 2023



Q1: Why not make a function which creates a child and replace the image itself?

Q2: why not use ~~try catch~~ for error handling?

Date 30th Oct, 2023

SP

FILE HANDLING :-

I/O Devices are treated as files.

cp src dst

CP f1.txt f2.txt

Copies contents of f1 into f2

1. open source file for (reading)
2. Open dst file for (writing). / created by file
3. Read from src
4. Write to dst
5. Close all

int open(~~const~~ char* path, int oflags, ...)

error :-1

success: > 0

open flags

optional

arg

int fd

int main () {

int perm = S_IRWXU | S_IRGRP | S_IWGRP | S_IROTH;

int fd1 = open("f1.txt", O_RDONLY);

permission

// error checking B must

if (fd1 == -1) { perror("Failed to open
f1 for reading (in)"); return 1; }

// int fd2 = open("f2.txt", O_WRONLY);

starts by writing,
erase full content

End mei writing
start karo.

oflags

O_RDONLY }
O_WRONLY }
O_RDWR }

O_CREAT }
O_EXCL }

error.
if file does not exist.

if already exist.

Mutually Exclusive → at a time one can be used.
Permissions.

S_IRUSR → for reading

S_IWUSR

S_IXUSR

S_IRWXU

Bitwise OR if f2. does not exist
O_WRONLY | O_CREAT, S_IRWXU | S_IRGRP |
S_IROTH);

O_CREAT must have third argument
for permission.

// CP will overwrite f2.

int fd2 = open("f2.txt", O_WRONLY | O_CREAT,
O_TRUNC)

// error checking.

// O_EXCL, O_TRUNC and O_APPEND are
also mutually exclusive.

// Read from Src.

char buff[100];

// Here comes read() to the
rescue.

// int read(int fd, char *buff,
int size of buffer);

This is my string

bytes returned
int br = read(fd1, buff, 100);
read // error checking.
success: 0
0 > 0

Size of b

If we used read again, then it will return 0.
→ End of file / Empty file.

// Now write

|| int write (int fd, const char *buff
(
int length of buff).

int bw = write (fd2, buff, br);

if 80 100 is written
then it will write garbage values.

close (fd1);
close (fd2);

}

*

Date 6 / 11 / 2023

File Representation :-

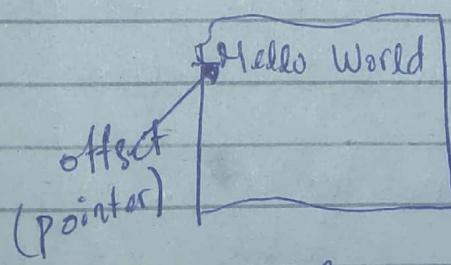
```
int main() {
```

```
    int fd1 = open ("f1.txt", O_RDONLY);  
    = read(fd1, buff, 5);
```

user space

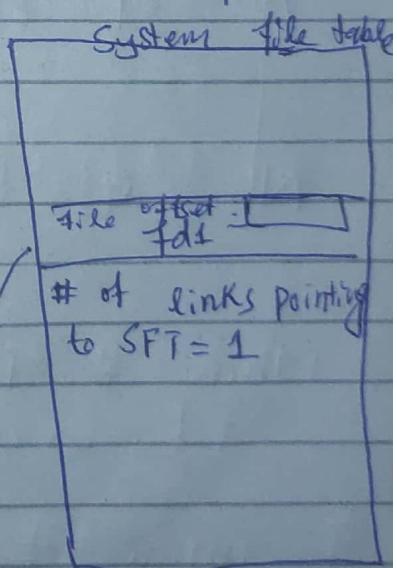
fd table

0	IN
1	OUT
2	ERR
3	fd1



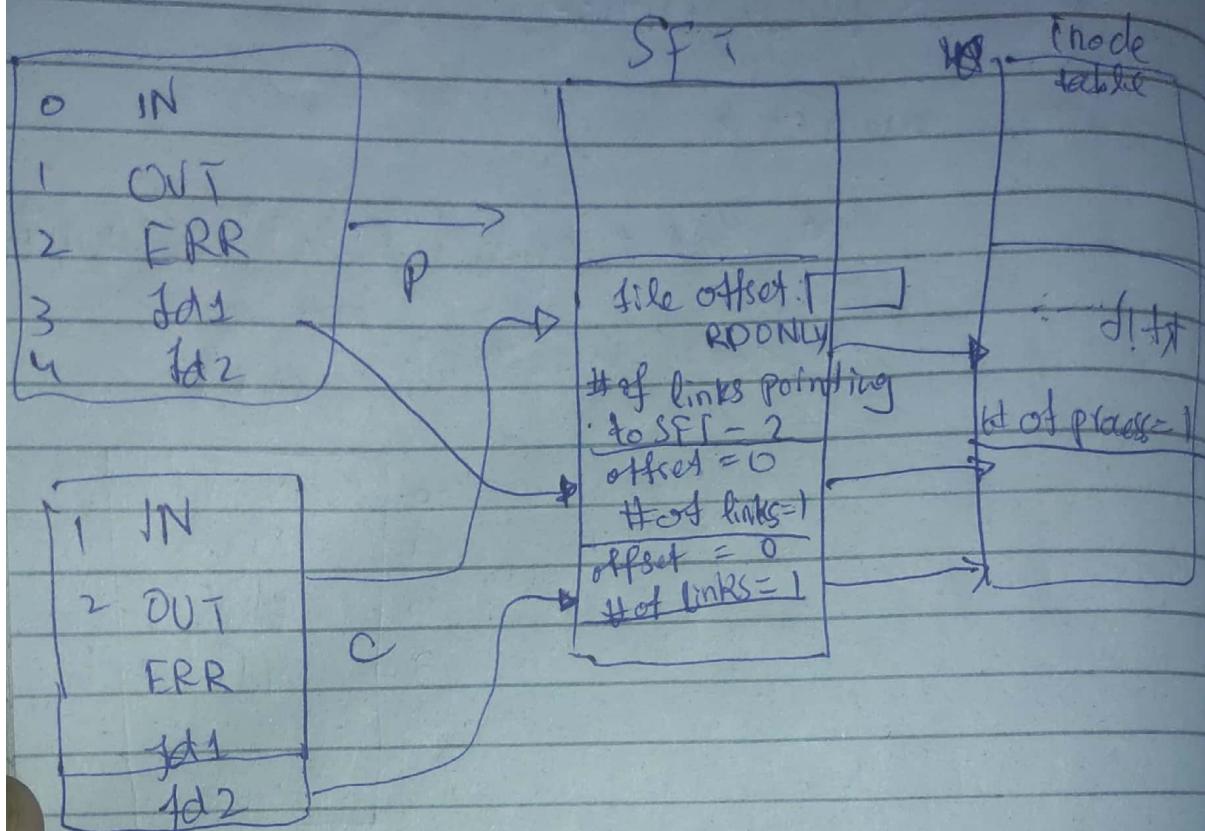
inode table

Kernel Space



shared b/w
all processes

```
/home/user/Desktop/  
f1.txt  
# of process: 1
```



No. of Entries in SFT = No. of times open is called

No. of Entries in Inode table = No. of actual files

* When # of links == 0, then SFT entry will be deleted.

* When # of process == 0, then Inode table entry will be deleted.

```

int main() {
    char buff[10];
    offset is
    shared
    int fd1 = open("f1.txt", O_RDONLY);
    int n = fork();
    if (n == 0) {
        int fd2 = open("f2.txt", O_WRONLY);
        int br = read(fd2, buff, 5);
        printf(buff);
    }
    else {
        int fd2 = open("f2.txt", O_WRONLY);
        int br = read(fd2, " ", 1);
        printf(" ");
    }
}

```

Filters & Redirection :-

① Cat → No Redirection
 Will read from STDIN, write to STDOUT by default

IN
f1.txt
STDERR

② Cat f1.txt → Input Redirection
 read from f1.txt, write to STDOUT

IN
OUT
ERR

③ $\text{cat} > \text{f1.txt}$

↳ read from **STDIN**, write to **f1.txt**
↳ output redirection

④

$\text{cat } \text{f1.txt} > \text{f2.txt}$

↳ IN/OUT redirection
Read from **f1.txt**, write to **f2.txt**.

`int dup2(int fd1, int fd2);`

for changing fd entries.

IN
OUT
ERR
<u>fd 1.txt</u>

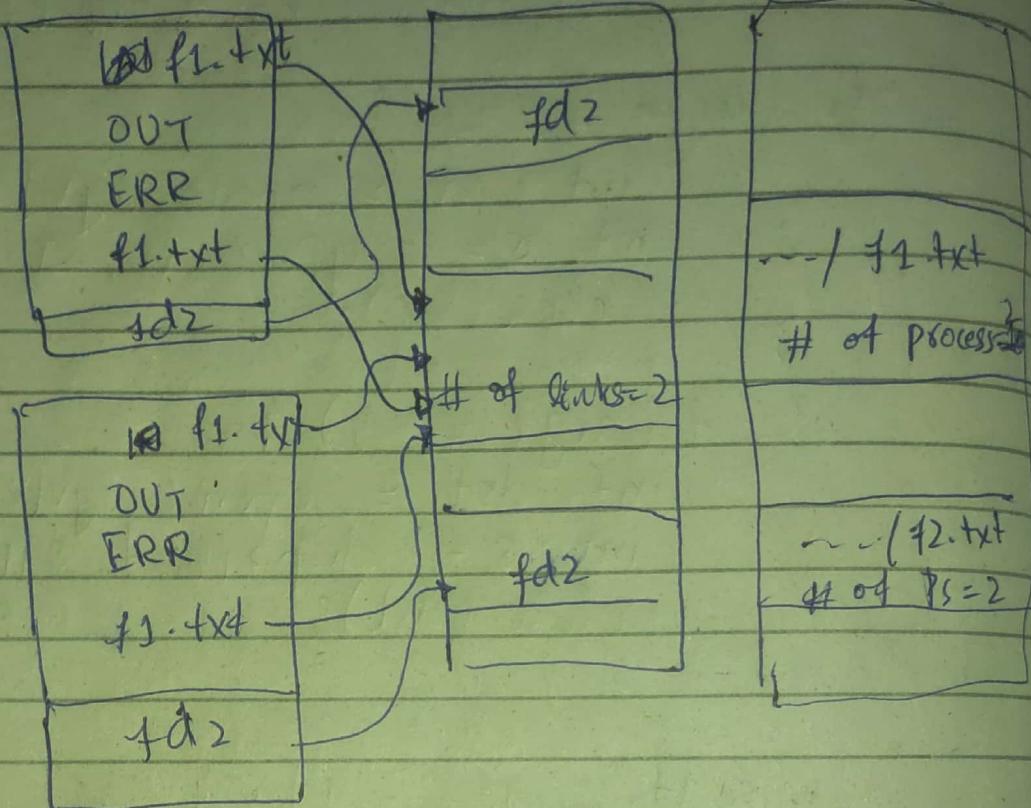
`dup2(fd1, STDIN_FILENO);`

`close(fd1);`

fd	{ 0 }	fd	<u>f1.txt</u>
			OUT
			ERR
			<u>fd1.txt</u>

```
int main( int argc, char *argv[7] ) {
    if (argc == 2) {
        int src = open(argv[1], O_RDONLY);
        int s = dup2(src, STDIN_FILENO);
        close(src);
    }
    else if (argc == 3) {
        int dst = open(argv[2], O_WRONLY);
        int s = dup2(dst, STDOUT_FILENO);
        close(dst);
    }
    else if (argc == 4) {
        int src = open(argv[1], O_RDONLY);
        int dst = open(argv[3], O_WRONLY);
        dup2(src, STDIN_FILENO);
        dup2(dst, STDOUT_FILENO);
        close(src);
        close(dst);
    }
    read(STDIN_FILENO
    write
```

SFT



```
int main( int argc, char *argv[] )
```

```
    int fd = open("f1.txt", O_RDONLY);
    dup2(fd, STDIN_FILENO);
    int x = fork();
    int fd2 = open("f2.txt", O_RDONLY);
    if (x == 0)
        close(fd);
```

3